Investigating Overfitting in Convolutional Neural Networks

1. Introduction

Neural networks are a form of machine learning models, roughly inspired by how the human brain processes information. They have become increasingly prevalent in our developing world, providing powerful and unique solutions to a wide variety of problems. However, with power also comes risk, as training a neural network too much on a small dataset can cause overfitting. Overfitting occurs when a model learns the training data too well by essentially memorizing it, causing it to perform well on data in its training set but poorly on unseen data. While it might seem that simply increasing a model's size would always boost accuracy by allowing it to capture more patterns, extra parameters often give the network enough capacity to store the training set outright, allowing it to overfit. Smaller networks, by contrast, are forced to find the underlying patterns in the data because of their limited size. In this study, we investigate whether increasing the number of parameters in a convolutional neural network leads to greater overfitting when trained on a small dataset, as measured by the difference between training and test accuracy.

2. Statistical Question

Does increasing the number of parameters in a fixed CNN architecture cause a greater difference between training and testing accuracy?

Hypotheses:

$$H_0 : \beta = 0$$

$$H_a : \beta > 0$$

Where β is the true slope of the population least-squares regression line that relates number of parameters of the model to the difference in accuracy of the model on the train dataset and the test dataset.

3. Data Collection

We trained $\mathbf{x}$ convolutional models on a randomly selected small subset of size 100 images from the CIFAR-10 dataset, each with the same architecture (*see figure 1*). Each model had a different number of parameters, roughly uniformly ranging from approximately 1 million to 50 million. We controlled the parameter count by multiplying the amount of filters for the Conv2D layers and the amount of units for the hidden Dense layers by a scalar factor, n. For each model, the initial values of the parameters are randomly set, so each model can be seen as randomly selected from all possible models of that size and architecture before training.

Model Architecture

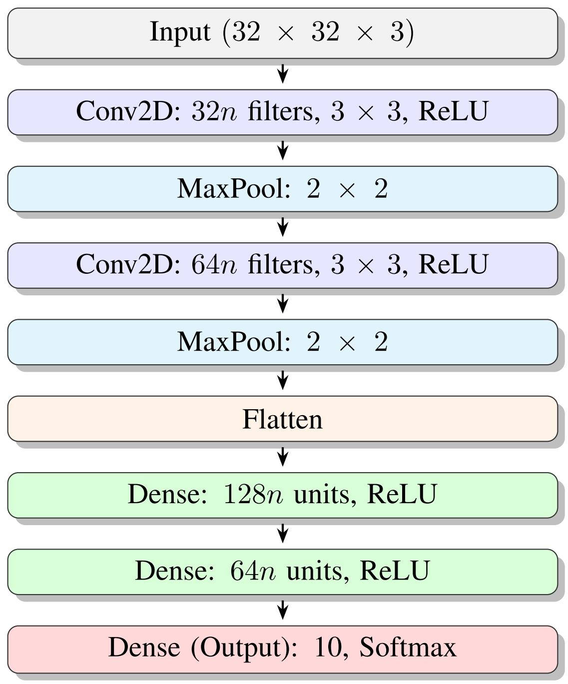

Figure 1: General Architecture of Our Models

Data Display

Data Analysis

Conclusion

Reflection